CSCE 580: Introduction to AI

Week 6 - Lectures 11 and 12:
AI Trust; Symbolic - Representation and Logic

PROF. BIPLAV SRIVASTAVA, AI INSTITUTE

23RD AND 25TH SEP 2025

Carolinian Creed: "I will practice personal and academic integrity."

Credits: Copyrights of all material reused acknowledged

Organization of Week 6 - Lectures 11, 12

- Introduction Section
  - Recap from Week 5 (Lectures 9 and 10)
  - AI news
- Main Section
  - L11: ML Trust Issues – Explainability, Rating
  - L12: Symbolic - Representation and Logic
- Concluding Section
  - About next week – W7: Lectures 13, 14
  - Ask me anything

Relationship Between Main AI Topics (Covered in Course)

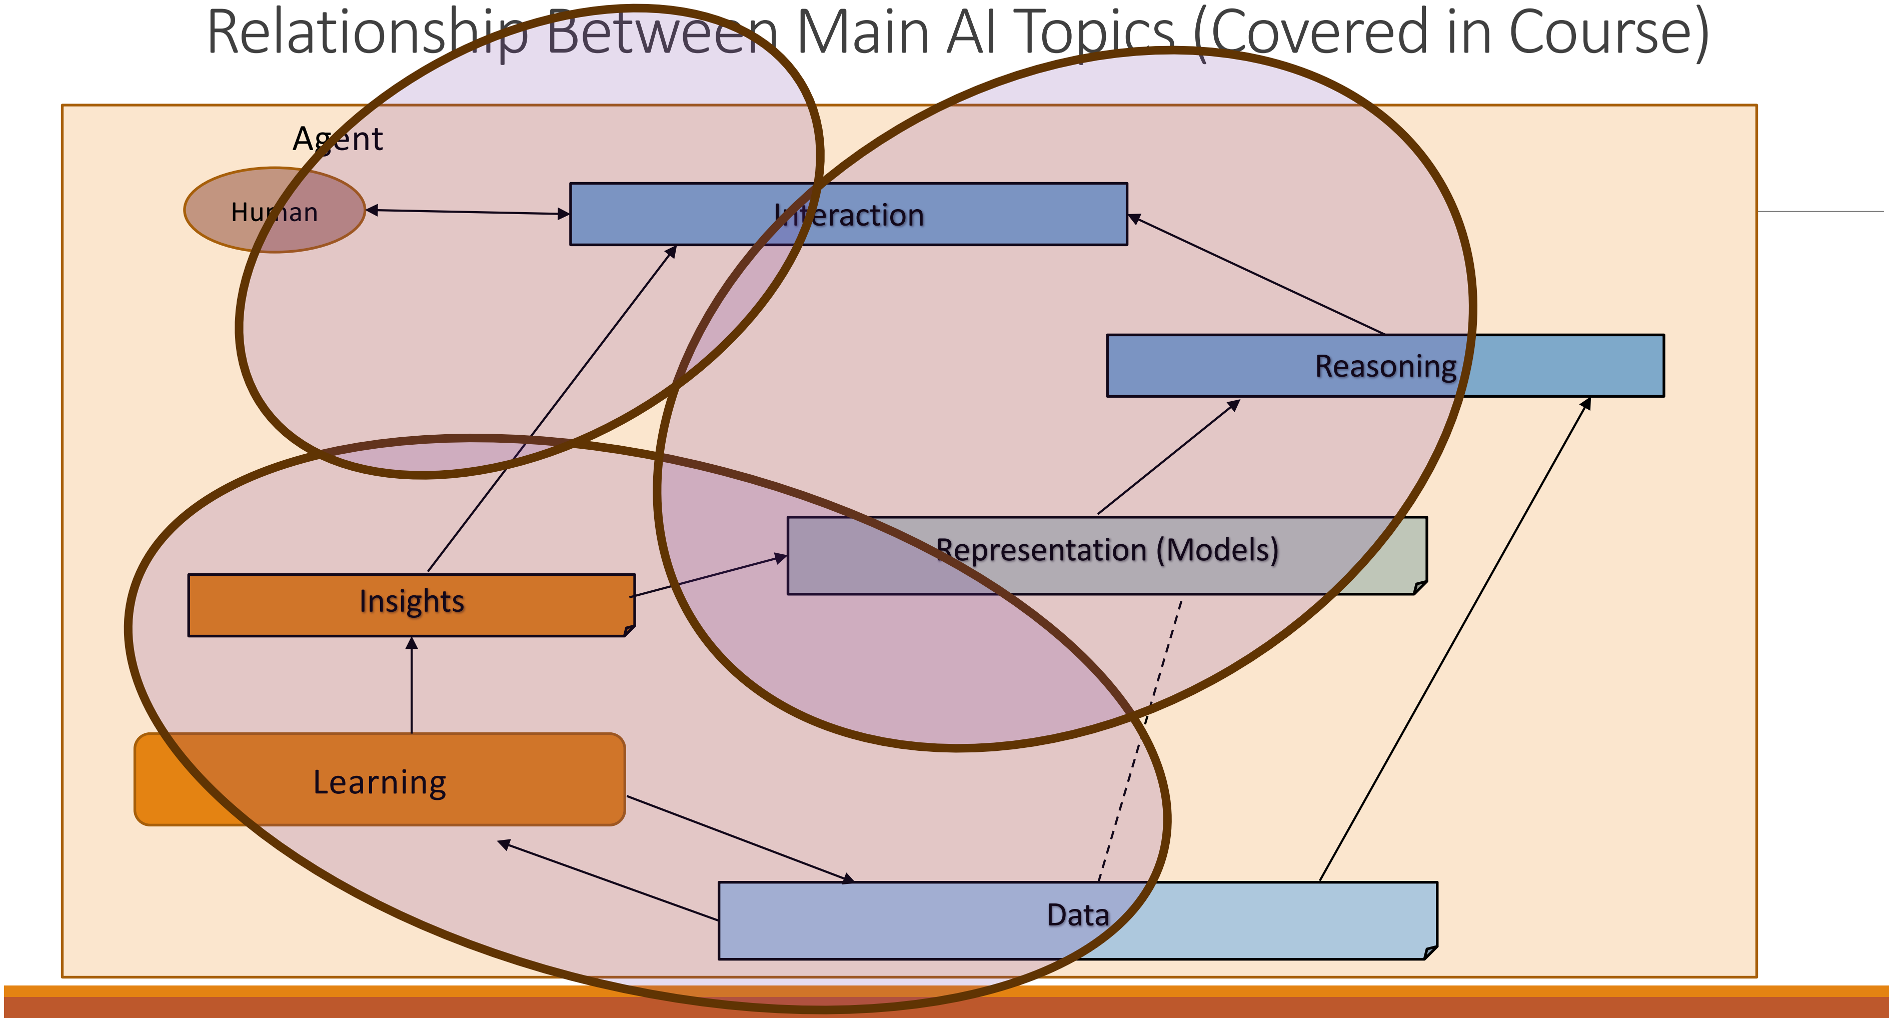

Recap of Week 5

- We talked about
  - Language models
  - LLMs
  - Using LLMs
  - AI/ LLM Trust
  - Exercise 1, Project A

- Week 1: Introduction, Aim: Chatbot / Intelligence Agent
- Weeks 2-3: Data: Formats, Representation and the Trust Problem
- Week 3: Machine Learning – Supervised (Classification)
- Week 4: Machine Learning - Unsupervised (Clustering) –
- Topic 5: Learning neural network, deep learning, Adversarial attacks
- Week 6: Large Language Models – Representation and Usage issues
- Weeks 7-8: Search, Heuristics - Decision Making
- Week 9: Constraints, Optimization – Decision Making
- Topic 10: Markov Decision Processes, Hidden Markov models - Decision making
- Topic 11-12: Planning, Reinforcement Learning – Sequential decision making
- Week 13: Trustworthy Decision Making: Explanation, AI testing
- Week 14: AI for Real World: Tools, Emerging Standards and Laws; Safe AI/ Chatbots

AI News

#1 NEWS – To fill

- Report: https://www.anthropic.com/news/detecting-countering-misuse-aug-2025

Press: https://www.nbcnews.com/tech/security/hacker-used-ai-automate-unprecedented-cybercrime-spree-anthropic-says-rcna227309

Key points

- “**first publicly documented instance** in which a hacker used a leading AI company’s chatbot to automate almost an entire cybercrime spree.”
- “(used Claude) to research, hack and extort at least 17 companies..”

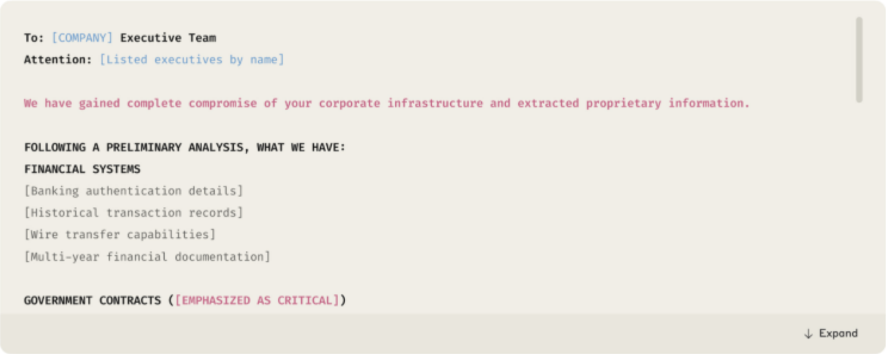

```
To: [COMPANY] Executive Team
Attention: [Listed executives by name]

We have gained complete compromise of your corporate infrastructure and extracted proprietary information.

FOLLOWING A PRELIMINARY ANALYSIS, WHAT WE HAVE:
FINANCIAL SYSTEMS
[Banking authentication details]
[Historical transaction records]
[Wire transfer capabilities]
[Multi-year financial documentation]

GOVERNMENT CONTRACTS ([EMPHASIZED AS CRITICAL])
```

A simulated custom ransom note. This is an illustrative example, created by our threat intelligence team for research and demonstration purposes after our analysis of extracted files from the real operation.

1. Specializes in “vibe coding,” or creating computer programming based on simple requests — to identify companies vulnerable to attack.
2. Claude then created malicious software to actually steal sensitive information from the companies.
3. Next, it organized the hacked files and analyzed them to both help determine what was sensitive and could be used to extort the victim companies.
4. Analyzed the companies’ hacked financial documents to help determine a realistic amount of bitcoin to demand in exchange for the hacker’s promise not to publish that material.
5. It also wrote suggested extortion emails.

Introduction Section

Lecture 11: Overcoming ML Trust Issues – Explainability, Rating

Lecture 11: Concluding Comments

We discussed

- AI trust and risk mitigation
- Explanability methods
- Rating methods

Lecture 12: Representation and Logic

Main Section

Intelligent Agent – Simple Knowledge Based

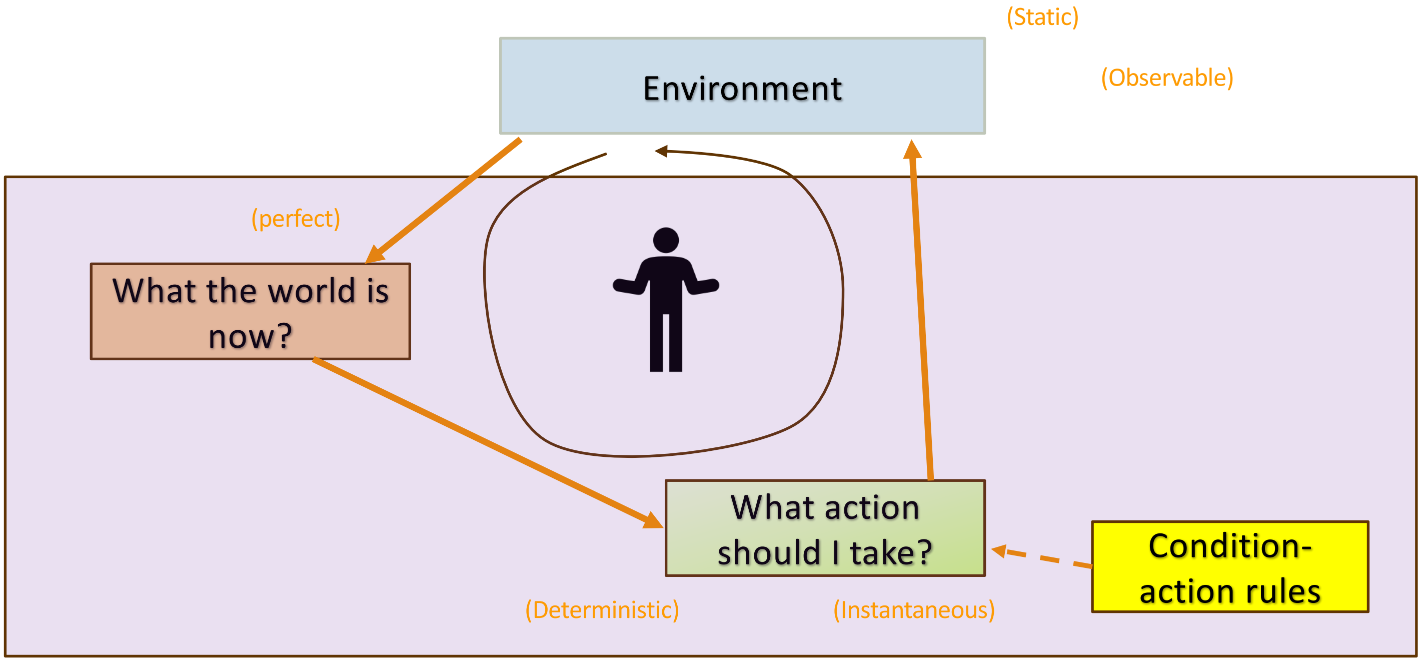

(Static)

Environment

(Observable)

(perfect)

What the world is now?

What action should I take?

Condition-action rules

(Deterministic)

(Instantaneous)

Logic – Basic Idea

- Starting with true assumptions, a knowledge-based system (automaton) to draw true conclusions
- Logic consists of three components
  - Syntax -- allowable sentences
  - Semantics -- determining truth of sentence given a model (assignment of values to sentences)
  - Inference Procedure -- rules to draw conclusion given a set of sentences

Formal Logic – 1/3

- An automaton for manipulating symbols and drawing conclusions
- Consists of a knowledge base with:
  - a set of true statements (sentences). Sentences have
    - Syntax
    - Semantics – compositional property
  - Proof theory: a set of rules for deducing the entailments / interpretations of the sentences

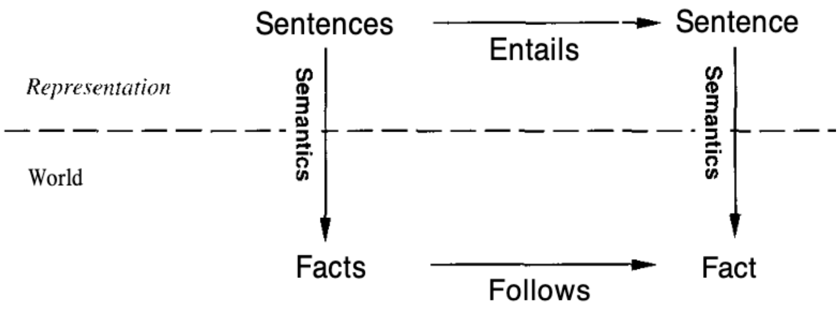

- Properties of sentences
  - **Valid**: A sentence is **valid** or necessarily true if and only if it is true under all possible interpretations in all possible worlds. Also called a **tautology**
  - **Satisfiable:** A sentence is satisfiable if and only if there is some interpretations in some possible worlds where it is true.

Credits:

- Russell & Norvig, AI - A Modern Approach
- Deepak Khemani - A First Course in AI

Propositional Logic

- Sentences: assertions about the world

  - Atomic sentence: propositional symbol
    - ClassToday – whether there is a class today: Yes, No
    - A -- any fact of interest: Yes, No
    - True -- always true
    - False -- always false

Propositional Logic - Syntax

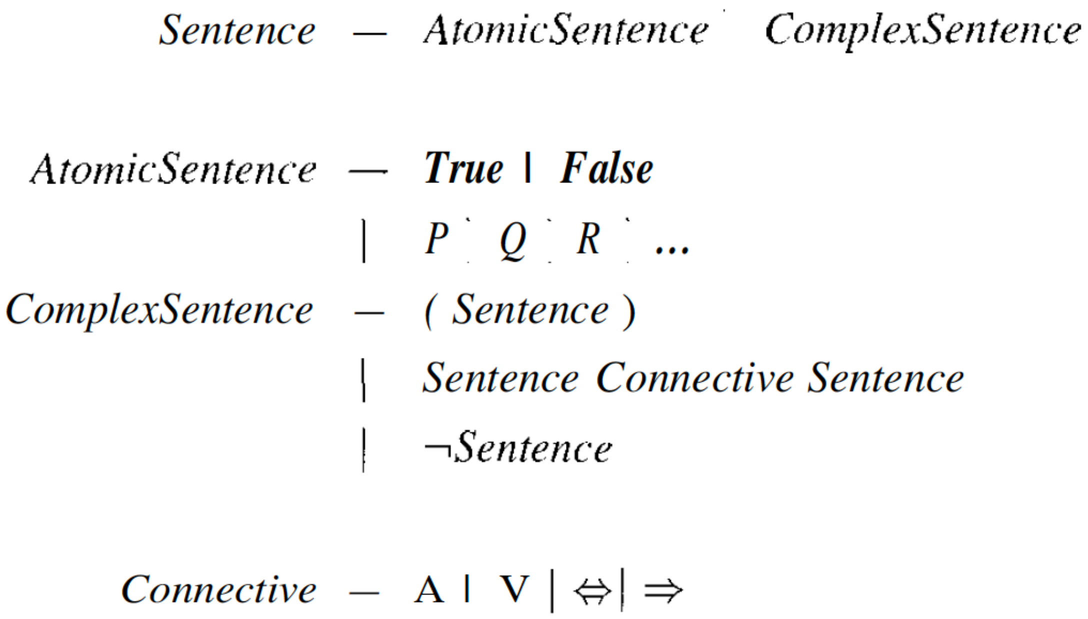

Sentence — *AtomicSentence* *ComplexSentence*

AtomicSentence — ***True*** | ***False***

| *P* *Q* *R* ...

ComplexSentence — (*Sentence*)

| *Sentence Connective Sentence*

| ¬*Sentence*

Connective — A | V | ⇔| ⇒

BNF (Backus-Naur Form) grammar of sentences in propositional logic

Source: Russell & Norvig, AI: A Modern Approach

Propositional Logic - Semantics

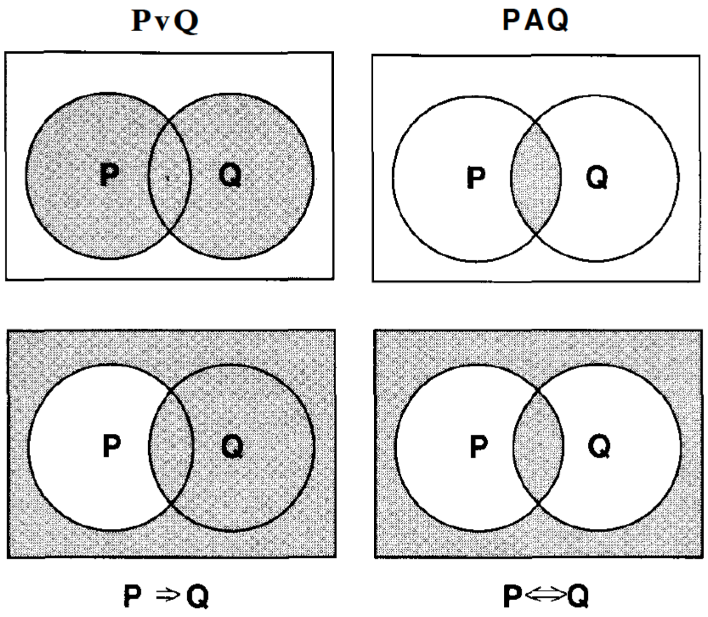

PvQ

PAQ

P ⇒Q

P⇔Q

Model of sentence: Any world in which a sentence is true (under a particular interpretation)

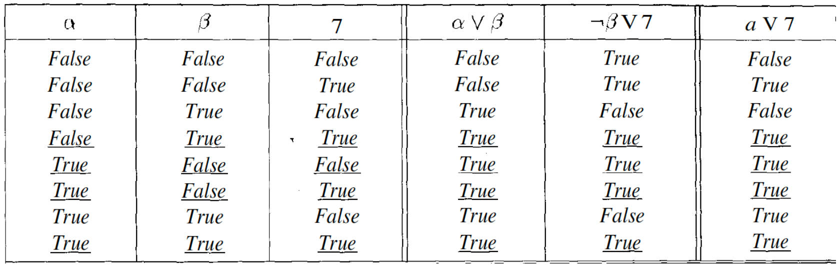

α	β	7	α ∨ β	¬β V 7	a V 7
False	*False*	*False*	*False*	*True*	*False*
False	*False*	*True*	*False*	*True*	*True*
False	*True*	*False*	*True*	*False*	*False*
False	*True*	*True*	*True*	*True*	*True*
True	*False*	*False*	*True*	*True*	*True*
True	*False*	*True*	*True*	*True*	*True*
True	*True*	*False*	*True*	*False*	*True*
True	*True*	*True*	*True*	*True*	*True*

Truth table to prove soundness of inference

Source: Russell & Norvig, AI: A Modern Approach

Inference Procedure

- Given a knowledge base (KB), generate new sentences α that are entailed by KB
  - KB |= α

- Given KB and α, report whether or not α is entailed by KB
  - KB |- α

-

Propositional Logic

Inference Procedures

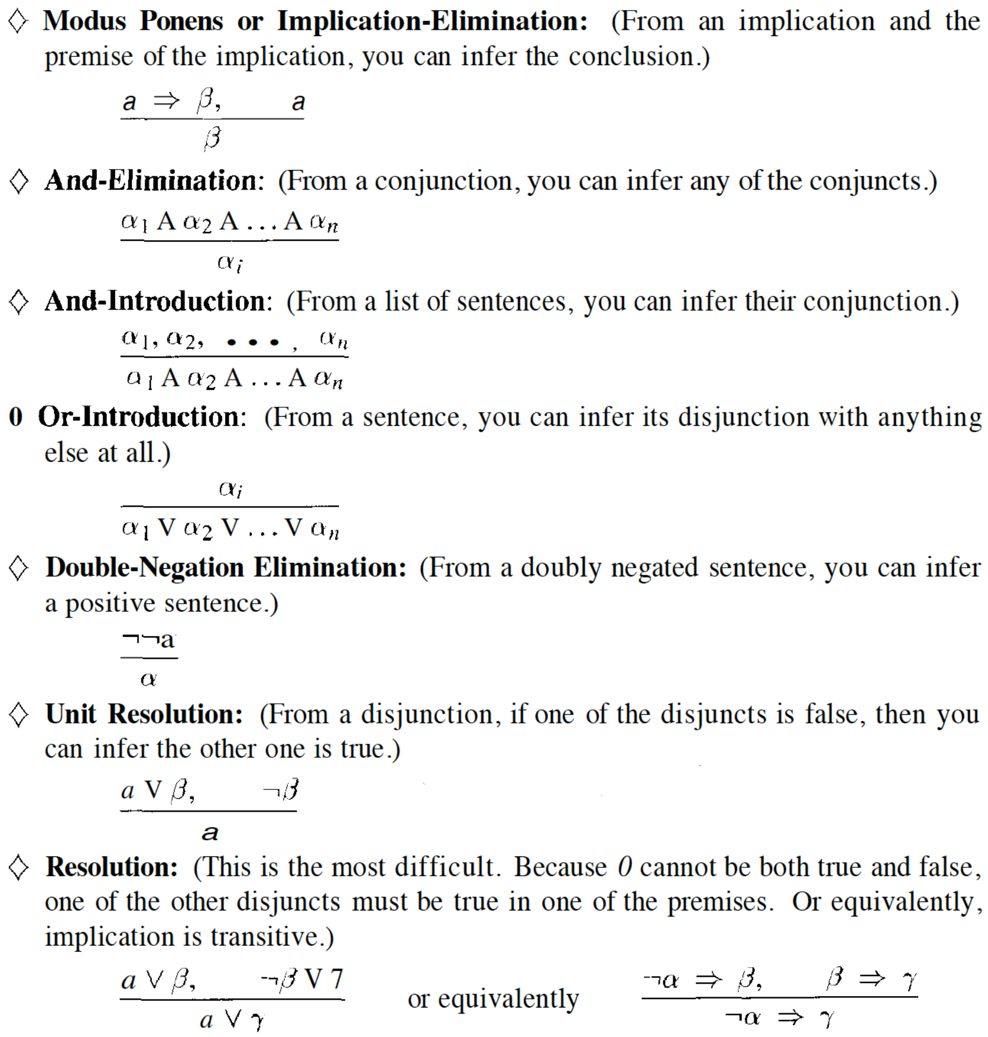

◇ **Modus Ponens or Implication-Elimination:** (From an implication and the premise of the implication, you can infer the conclusion.)

$$\frac{a \Rightarrow \beta, \qquad a}{\beta}$$

◇ **And-Elimination**: (From a conjunction, you can infer any of the conjuncts.)

$$\frac{\alpha_1 \mathrm{A} \alpha_2 \mathrm{A} \ldots \mathrm{A} \alpha_n}{\alpha_i}$$

◇ **And-Introduction**: (From a list of sentences, you can infer their conjunction.)

$$\frac{\alpha_1, \alpha_2, \bullet\bullet\bullet, \alpha_n}{\alpha_1 \mathrm{A} \alpha_2 \mathrm{A} \ldots \mathrm{A} \alpha_n}$$

0 Or-Introduction: (From a sentence, you can infer its disjunction with anything else at all.)

$$\frac{\alpha_i}{\alpha_1 \mathrm{V} \alpha_2 \mathrm{V} \ldots \mathrm{V} \alpha_n}$$

◇ **Double-Negation Elimination:** (From a doubly negated sentence, you can infer a positive sentence.)

$$\frac{\neg\neg a}{\alpha}$$

◇ **Unit Resolution:** (From a disjunction, if one of the disjuncts is false, then you can infer the other one is true.)

$$\frac{a \mathrm{V} \beta, \qquad \neg\beta}{a}$$

◇ **Resolution:** (This is the most difficult. Because *0* cannot be both true and false, one of the other disjuncts must be true in one of the premises. Or equivalently, implication is transitive.)

$$\frac{a \vee \beta, \qquad \neg\beta \mathrm{V} 7}{a \vee \gamma} \qquad \text{or equivalently} \qquad \frac{\neg\alpha \Rightarrow \beta, \qquad \beta \Rightarrow \gamma}{\neg\alpha \Rightarrow \gamma}$$

Source: Russell & Norvig, AI: A Modern Approach

KB Agent Procedure

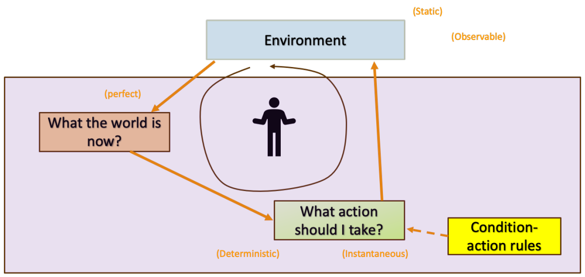

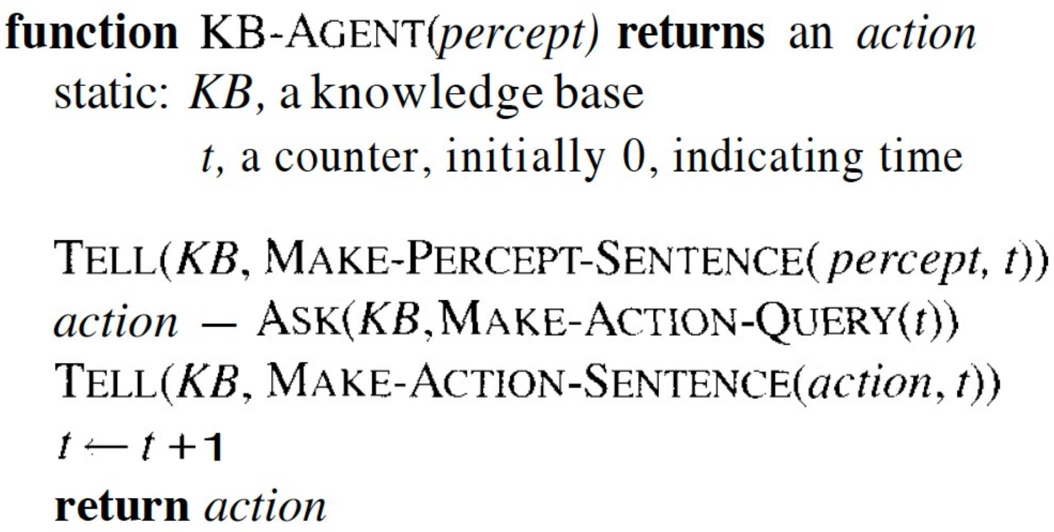

```
function KB-AGENT(percept) returns an action
    static: KB, a knowledge base
            t, a counter, initially 0, indicating time

    TELL(KB, MAKE-PERCEPT-SENTENCE(percept, t))
    action — ASK(KB, MAKE-ACTION-QUERY(t))
    TELL(KB, MAKE-ACTION-SENTENCE(action, t))
    t ← t + 1
    return action
```

Source: Russell & Norvig, AI: A Modern Approach

Formal Logic – 2/3

- Levels at which sentences are encoded
  - Epistemic (also called knowledge): what agents knows or believes
  - Logical: how sentences are encoded to allow inferencing. E.g., symbols
  - Executional: how sentences are encoded during execution. E.g., vectors, symbols

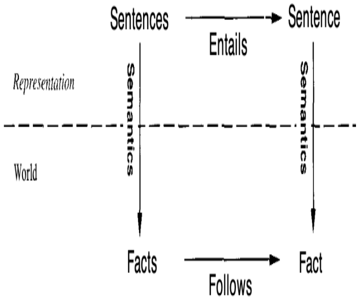

- Properties of sentences
  - **Valid**: A sentence is **valid** or necessarily true if and only if it is true under all possible interpretations in all possible worlds. Also called a **tautology**
  - **Satisfiable:** A sentence is satisfiable if and only if there is some interpretations in some possible worlds where it is true.

Credits:
- Russell & Norvig, AI - A Modern Approach
- Deepak Khemani - A First Course in AI

Formal Logic – 3/3

- Properties of Logic System
  - **Soundness**: if it produces only true statements
  - **Completeness**: if it produces all true statements
  - **Consistency**: if it does not produce a sentence and its negation

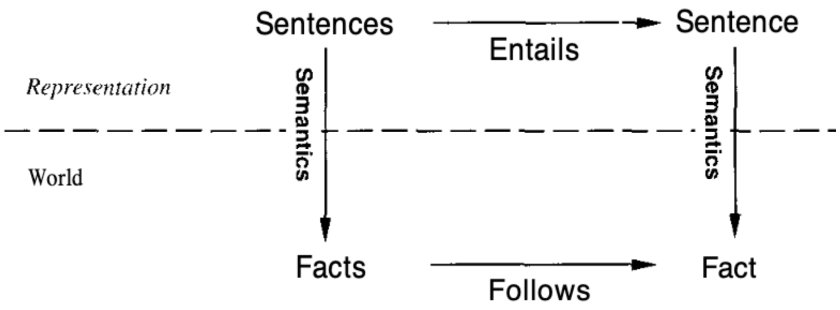

Sentences

Entails

Sentence

Representation

Semantics

Semantics

World

Facts

Follows

Fact

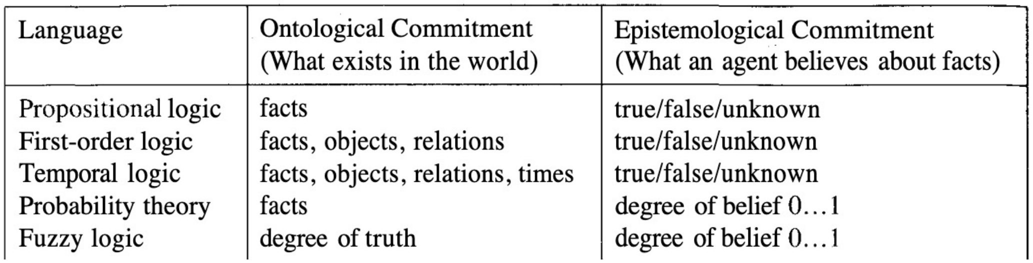

Language	Ontological Commitment (What exists in the world)	Epistemological Commitment (What an agent believes about facts)
Propositional logic	facts	true/false/unknown
First-order logic	facts, objects, relations	true/false/unknown
Temporal logic	facts, objects, relations, times	true/false/unknown
Probability theory	facts	degree of belief 0...1
Fuzzy logic	degree of truth	degree of belief 0...1

Credits:

- Russell & Norvig, AI - A Modern Approach
- Deepak Khemani - A First Course in AI

Example: Course Selection

Example Situation – Course Selection

- A person wants to pass an academic program in two majors: A and B
- There are three subjects available: A, B and C, each with three levels (*1, *2, *3). There are thus 9 courses: A1, A2, A3, B1, B2, B3, C1, C2, C3
- To graduate, at least one course at beginner (*1) level is needed in major(s) of choice(s), and two courses at intermediate levels (*2) are needed
- Answer questions
  - Q1: How many minimum courses does the person have to take ?
  - Q2: Can a person graduate in 2 majors studying 3 courses only?
  - …

Representation - Example

- Domain Description: “There are three subjects: A, B and C, each with three levels (*1, *2, *3).”
- Representation
  - has_studied_courseA1: yes – student has taken course; no – student has not taken
  - has_studied_courseA2
  - has_studied_courseA3
  - has_studied_courseB1
  - has_studied_courseB2
  - has_studied_courseB3
  - has_studied_courseC1
  - has_studied_courseC2
  - has_studied_courseC3

Issue: What about hierarchy among courses?

Representation - Example

- Domain Description: “There are three subjects: A, B and C, each with three levels (*1, *2, *3).”
- Representation
  - has_studied_courseA1: yes – student has taken course; no – student has not taken
  - has_studied_courseA2
  - has_studied_courseA3
  - has_studied_courseB1
  - has_studied_courseB2
  - has_studied_courseB3
  - has_studied_courseC1
  - has_studied_courseC2
  - has_studied_courseC3

LowerThan_Course_A1_CourseA2
LowerThan_Course_A2_CourseA3
LowerThan_Course_B1_CourseB2
LowerThan_Course_B2_CourseB3
LowerThan_Course_C1_CourseC2
LowerThan_Course_AC_CourseC3

- Previous statements set did not capture hierarchy between levels; new sentences would not have followed the reality in the world. Need more statements – LowerThan as shown.

How to Tackle Course Selection Problem ?

- Represent the world as sentences in KB
  - Update KB based on scenarios
- Solve problems about courses selection scenarios
  - Pose problems as queries to KB
  - Interpret answers // reason with the world

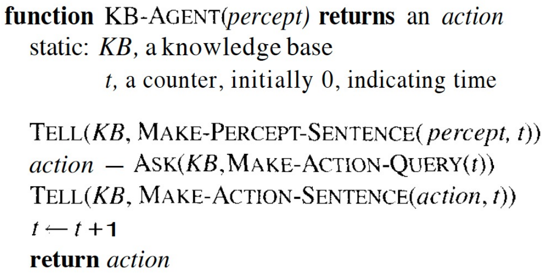

```
function KB-AGENT(percept) returns an action
  static: KB, a knowledge base
          t, a counter, initially 0, indicating time

  TELL(KB, MAKE-PERCEPT-SENTENCE(percept, t))
  action — ASK(KB,MAKE-ACTION-QUERY(t))
  TELL(KB, MAKE-ACTION-SENTENCE(action, t))
  t ← t + 1
  return action
```

Source: Russell & Norvig, AI: A Modern Approach

Major Types of Reasoning

- Inference: From premises to conclusions
  - Major types
    - **Deduction**: deriving logical conclusions from premises known or assumed to be true
    - **Induction**: deriving from particular premises to a universal conclusion.
    - **Abduction**: from an observation, find the most likely conclusion from the observations

- Usage
  - Deduction is useful to build knowledge bases from parts
  - Induction: to generalize
  - Abduction is a good source for hypothesis / priors in Bayesian learning

Setting Up for AIMA Code

- AI resources
  https://github.com/biplav-s/course-ai-tai-f23/blob/main/sample-code/AI-Resources.md
- Setting up for Python code from AIMA book
  https://github.com/biplav-s/course-ai-tai-f23/tree/main/sample-code/ai-book-samples

Exercise and Code

- Logical Reasoning
  - From Book: AI – A Modern Approach, https://github.com/aimacode/aima-python/blob/master/logic.ipynb

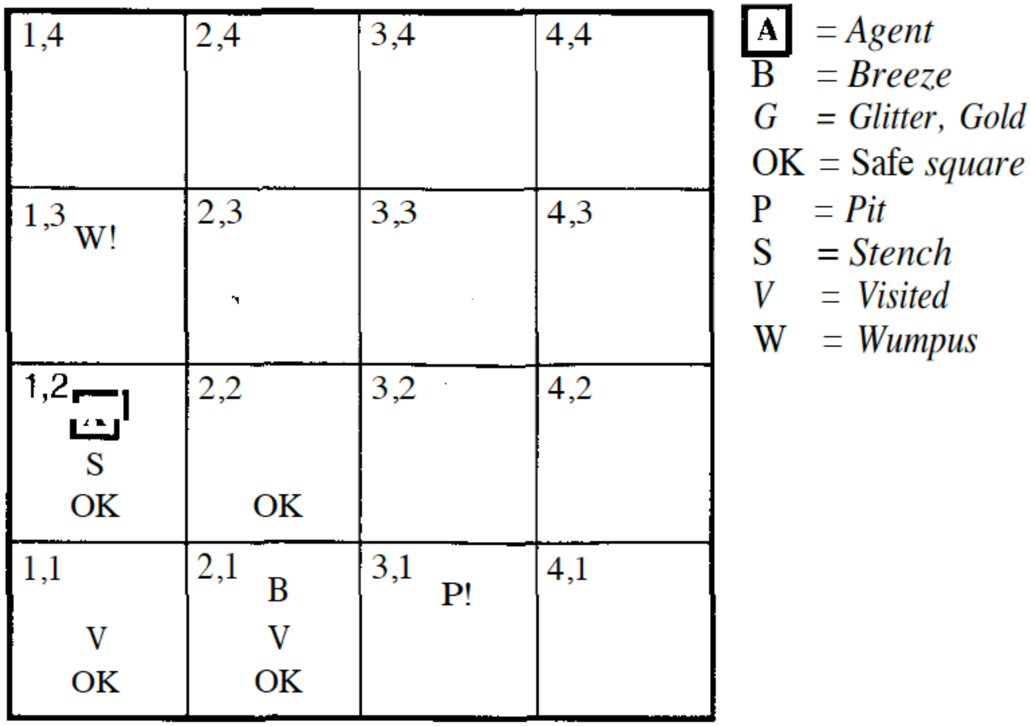

1,4	2,4	3,4	4,4
1,3 W!	2,3	3,3	4,3
1,2 A S OK	2,2 OK	3,2	4,2
1,1 V OK	2,1 B V OK	3,1 P!	4,1

A	= *Agent*
B	= *Breeze*
G	= *Glitter, Gold*
OK	= Safe *square*
P	= *Pit*
S	= *Stench*
V	= *Visited*
W	= *Wumpus*

Source: Russell & Norvig, AI: A Modern Approach

Examples of Agents

Agent Type	Percepts	Actions	Goals	Environment
Medical diagnosis system	Symptoms, findings, patient's answers	Questions, tests, treatments	Healthy patient, minimize costs	Patient, hospital
Satellite image analysis system	Pixels of varying intensity, color	Print a categorization of scene	Correct categorization	Images from orbiting satellite
Part-picking robot	Pixels of varying intensity	Pick up parts and sort into bins	Place parts in correct bins	Conveyor belt with parts
Refinery controller	Temperature, pressure readings	Open, close valves; adjust temperature	Maximize purity, yield, safety	Refinery
Interactive English tutor	Typed words	Print exercises, suggestions, corrections	Maximize student's score on test	Set of students

Source: Russell & Norvig, AI: A Modern Approach

Lecture 12: Summary

- We talked about
  - Knowledge-based agents
  - Logic (Propositional)
  - Inferencing (Propositional)

Week 6: Concluding Comments

We talked about

- AI/ ML Trust
  - Explainability
  - Trust ratings
- Representation and Logic
  - Propositional

- Week 1: Introduction, Aim: Chatbot / Intelligence Agent
- Weeks 2-3: Data: Formats, Representation and the Trust Problem
- Week 3: Machine Learning – Supervised (Classification)
- Week 4: Machine Learning - Unsupervised (Clustering) –
- Topic 5: Learning neural network, deep learning, Adversarial attacks
- Week 6: Large Language Models – Representation and Usage issues
- Weeks 7-8: Search, Heuristics - Decision Making
- Week 9: Constraints, Optimization – Decision Making
- Topic 10: Markov Decision Processes, Hidden Markov models - Decision making
- Topic 11-12: Planning, Reinforcement Learning – Sequential decision making
- Week 13: Trustworthy Decision Making: Explanation, AI testing
- Week 14: AI for Real World: Tools, Emerging Standards and Laws; Safe AI/ Chatbots

Upcoming Evaluation Milestones

- Projects B: Sep 30 – Nov 20
- Quiz 2: Oct 7
- Quiz 3: Nov 11
- Paper presentation (grad students only) : Nov 18

About Week 7 – Lectures 13 and 14

Week 7 – Lectures 13 and 14

- L13: Logic and Inference - First Order
- L12: Search, Search - Uninformed

- Week 1: Introduction, Aim: Chatbot / Intelligence Agent
- Weeks 2: Data: Formats, Representation, ML Basics
- Week 3: Machine Learning – Supervised (Classification)
- Week 4: Machine Learning - Unsupervised (Clustering) –
- Topic 5: Learning neural network, deep learning, Adversarial attacks
- Week 6: Large Language Models – Representation and Usage issues
- Weeks 7-8: Search, Heuristics - Decision Making
- Week 9: Constraints, Optimization – Decision Making
- Topic 10: Markov Decision Processes, Hidden Markov models - Decision making
- Topic 11-12: Planning, Reinforcement Learning – Sequential decision making
- Week 13: Trustworthy Decision Making: Explanation, AI testing
- Week 14: AI for Real World: Tools, Emerging Standards and Laws; Safe AI/ Chatbots

Note: exact schedule changes slightly to accommodate for exams and holidays.